

Assignment_8_2_Raghuwanshi_Prashant_DSC550

October 23, 2021

Assignment: 8.2 Exercise: Model Evaluation and Selection

Name: Prashant Raghuwanshi

Date: 10/23/2021

Course: DSC550-T301 Data Mining (2221-1)

Case Study: Titanic Case Study Part 3

```
[1]: import pandas as pd
import yellowbrick
import matplotlib.pyplot as plt
```

```
C:\ProgramData\Anaconda3\lib\site-packages\sklearn\utils\deprecation.py:144:
FutureWarning: The sklearn.metrics.classification module is deprecated in
version 0.22 and will be removed in version 0.24. The corresponding classes /
functions should instead be imported from sklearn.metrics. Anything that cannot
be imported from sklearn.metrics is now part of the private API.
    warnings.warn(message, FutureWarning)
```

```
[2]: # 1.Load the data from the "train.csv" file into a DataFrame.
addr1 = "C:/Users/dell/Documents/Machine_learning_assignments/week-6/train.csv"
df_train = pd.read_csv(addr1)
```

```
[3]: # 2.Display the dimensions of the file (so you'll have a good idea the amount
    ↳ of data you are working with.
print("The dimension of the table is: ", df_train.shape)
```

The dimension of the table is: (891, 12)

```
[4]: # 3.Display the first 5 rows of data so you can see the column headings and the
    ↳ type of data for each column.
print(df_train.head(5))
# a.Notice that Survived is represented as a 1 or 0
# b.Notice that missing data is represented as "NaN"
# c.The Survived variable will be the "target" and the other variables will be
    ↳ the "features"
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[5]: #4.      Think about some questions that might help you predict who will
      ↪survive:
      # a.What do the variables look like? For example, are they numerical or
      ↪categorical data. If they are numerical, what are their distribution; if
      ↪they are categorical, how many are they in different categories?
      # b.Are the numerical variables correlated?
      # c.Are the distributions of numerical variables the same or different among
      ↪survived and not survived? Is the survival rate different for different
      ↪values? For example, were people more likely to survive if they were younger?
      # d.Are there different survival rates in different categories? For example,
      ↪did more women survived than man?
```

```
[6]: # target variable is categorical variable, some of the feature variable are
      ↪numeric
```

```
[7]: #5.Look at summary information about your data (total, mean, min, max, freq,
      ↪unique, etc.)
print("\nDescribe Data\n")
print(df_train.describe())
```

Describe Data

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	714.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	14.526497	1.102743	

min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[8]: print("\nSummarized Data\n")
      print(df_train.describe(include=['O']))
```

Summarized Data

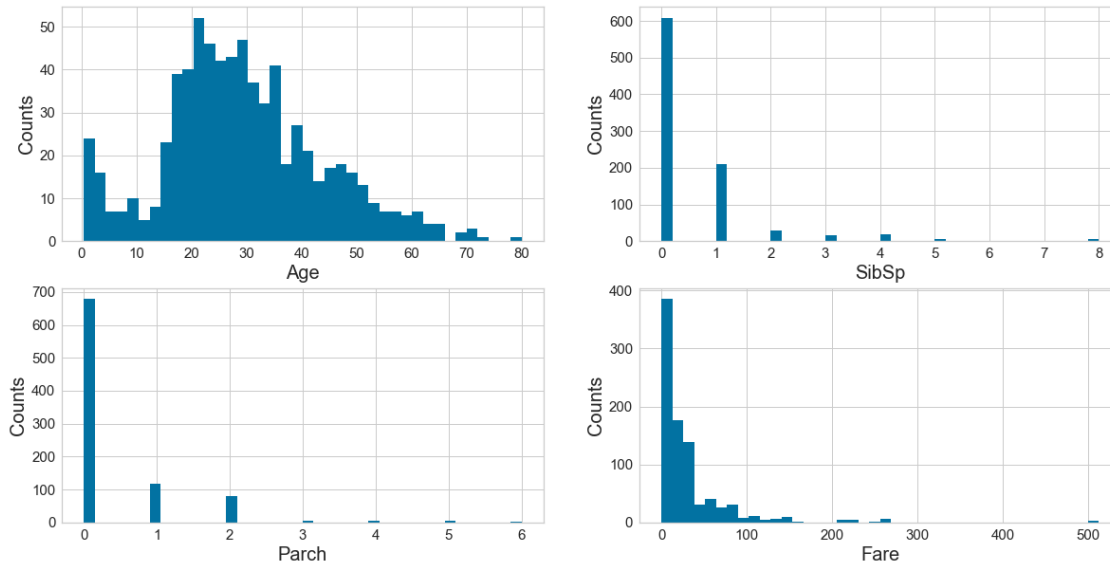
	Name	Sex	Ticket	Cabin	Embarked
count	891	891	891	204	889
unique	891	2	681	147	3
top	Vande Velde, Mr. Johannes Joseph	male	347082	C23 C25 C27	S
freq	1	577	7	4	644

```
[9]: #6.Make some histograms of your data ("A picture is worth a thousand words!")
      # a.Most of the passengers are around 20 to 30 years old and don't have
      ↳siblings or relatives with them.
      # b.A large amount of the tickets sold were less than $50. There are very few
      ↳tickets sold where the fare was over $500.
```

```
num_features = ['Age', 'SibSp', 'Parch', 'Fare']
xaxes = num_features
yaxes = ['Counts', 'Counts', 'Counts', 'Counts']
```

```
[10]: # set up the figure size
      plt.rcParams['figure.figsize'] = (20, 10)
      # make subplots
      fig, axes = plt.subplots(nrows = 2, ncols = 2)
      # draw histograms
      axes = axes.ravel()
      for idx, ax in enumerate(axes):
          ax.hist(df_train[num_features[idx]].dropna(), bins=40)
          ax.set_xlabel(xaxes[idx], fontsize=20)
          ax.set_ylabel(yaxes[idx], fontsize=20)
```

```
ax.tick_params(axis='both', labels=15)
plt.show()
```



```
[11]: #7.Make some bar charts for variables with only a few options.
# a.Ticket and Cabin have more than 100 variables so don't do those!
#%matplotlib inline
plt.rcParams['figure.figsize'] = (20, 10)
# make subplots
fig, axes = plt.subplots(nrows = 2, ncols = 2)
# make the data read to feed into the visulizer
X_Survived = df_train.replace({'Survived': {1: 'yes', 0: 'no'}}).
    ↳groupby('Survived').size().reset_index(name='Counts')['Survived']
Y_Survived = df_train.replace({'Survived': {1: 'yes', 0: 'no'}}).
    ↳groupby('Survived').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[0, 0].bar(X_Survived, Y_Survived)
axes[0, 0].set_title('Survived', fontsize=25)
axes[0, 0].set_ylabel('Counts', fontsize=20)
axes[0, 0].tick_params(axis='both', labels=15)
# make the data read to feed into the visulizer
X_Pclass = df_train.replace({'Pclass': {1: '1st', 2: '2nd', 3: '3rd'}}).
    ↳groupby('Pclass').size().reset_index(name='Counts')['Pclass']
Y_Pclass = df_train.replace({'Pclass': {1: '1st', 2: '2nd', 3: '3rd'}}).
    ↳groupby('Pclass').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[0, 1].bar(X_Pclass, Y_Pclass)
axes[0, 1].set_title('Pclass', fontsize=25)
axes[0, 1].set_ylabel('Counts', fontsize=20)
```

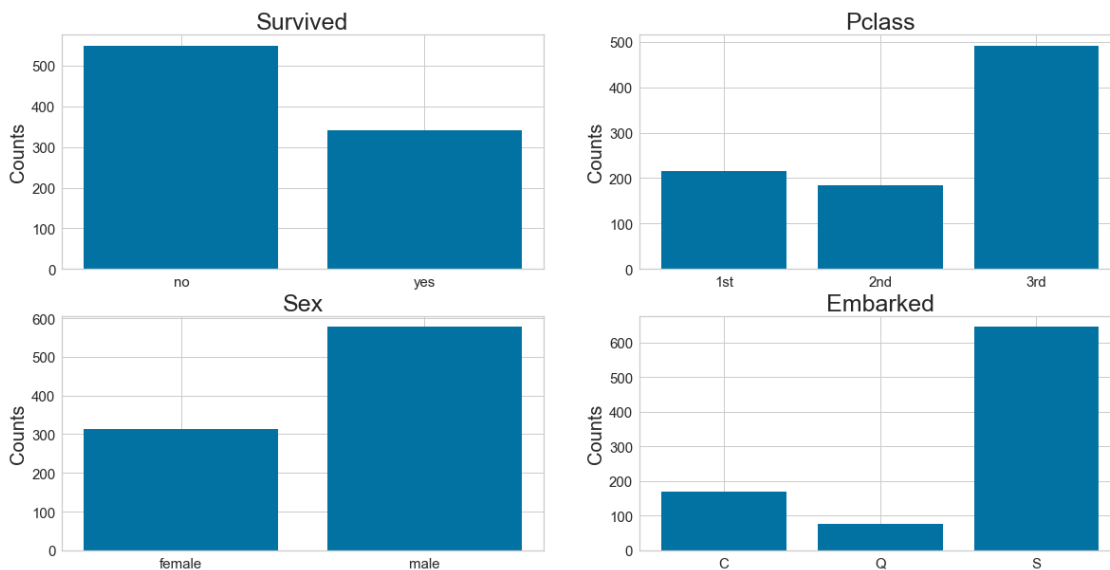
```

axes[0, 1].tick_params(axis='both', labels=15)

# make the data read to feed into the visulizer
X_Sex = df_train.groupby('Sex').size().reset_index(name='Counts')['Sex']
Y_Sex = df_train.groupby('Sex').size().reset_index(name='Counts')['Counts']
# make the bar plot
axes[1, 0].bar(X_Sex, Y_Sex)
axes[1, 0].set_title('Sex', fontsize=25)
axes[1, 0].set_ylabel('Counts', fontsize=20)
axes[1, 0].tick_params(axis='both', labels=15)

# make the data read to feed into the visulizer
X_Embarked = df_train.groupby('Embarked').size().
    ↪reset_index(name='Counts')['Embarked']
Y_Embarked = df_train.groupby('Embarked').size().
    ↪reset_index(name='Counts')['Counts']
# make the bar plot
axes[1, 1].bar(X_Embarked, Y_Embarked)
axes[1, 1].set_title('Embarked', fontsize=25)
axes[1, 1].set_ylabel('Counts', fontsize=20)
axes[1, 1].tick_params(axis='both', labels=15)
#plt.show()

```



[12]: #8.To see if the data is correlated, make some Pearson Ranking charts
 # b.The correlation between the variables is low (1 or -1 is high positive or ↪
 ↪high negative, 0 is low or no correlation)
 # These results show there is "some" positive correlation but it's not a ↪
 ↪high correlation.

```

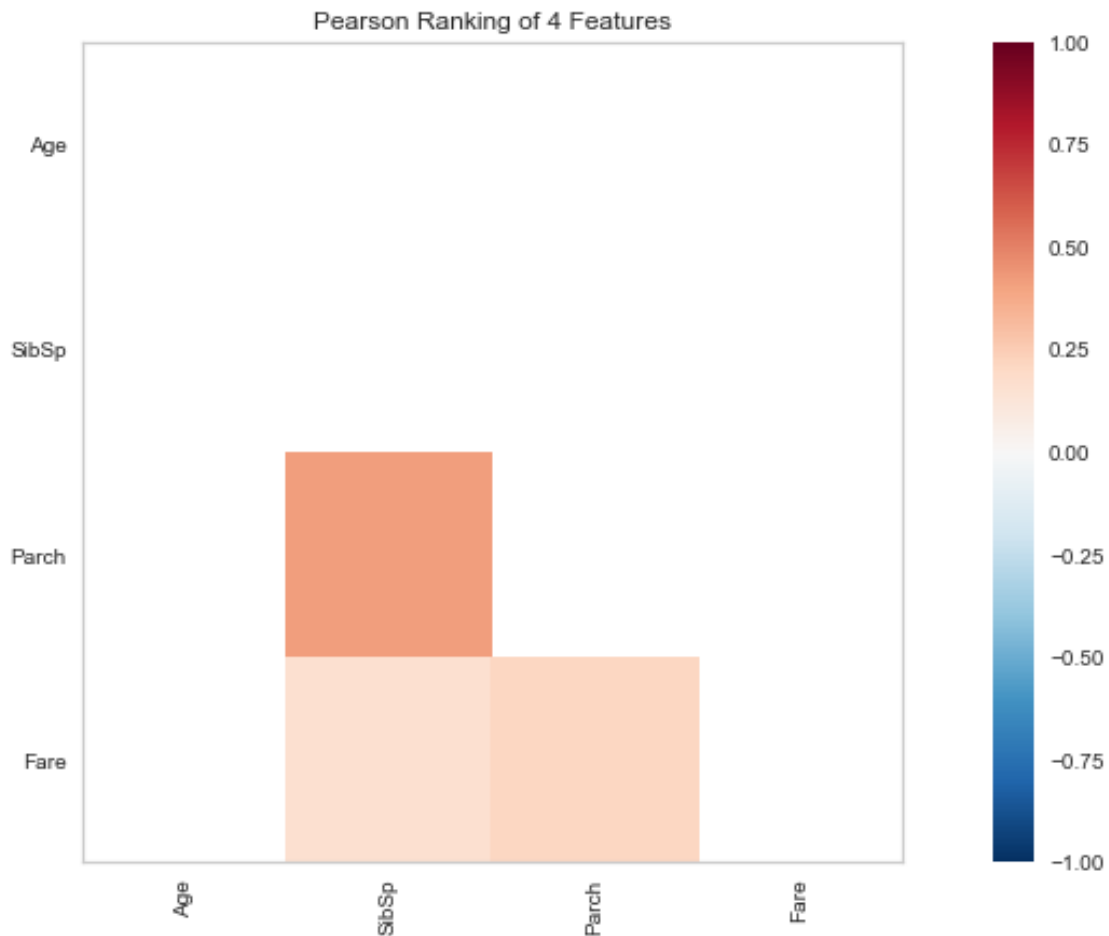
plt.rcParams['figure.figsize'] = (15, 7)

# import the package for visulization of the correlation
from yellowbrick.features import Rank2D

# extract the numpy arrays from the data frame
X = df_train[num_features].to_numpy()

# instantiate the visualizer with the Covariance ranking algorithm
visualizer = Rank2D(features=num_features, algorithm='pearson')
visualizer.fit(X) # Fit the data to the visualizer
visualizer.transform(X) # Transform the data
# a.Notice that in my sample code, I have saved this png file.
visualizer.poof(outpath="d://pcoords1.png") # Draw/show/poof the data
#plt.show()

```



```
[13]: #Use Parallel Coordinates visualization to compare the distributions of
      ↪numerical variables between passengers that survived and those that did not
      ↪survive.
      #a. That's a cool chart, isn't it?!
      # Compare variables against Survived and Not Survived
      #set up the figure size
      #%matplotlib inline
      plt.rcParams['figure.figsize'] = (15, 7)
      plt.rcParams['font.size'] = 50

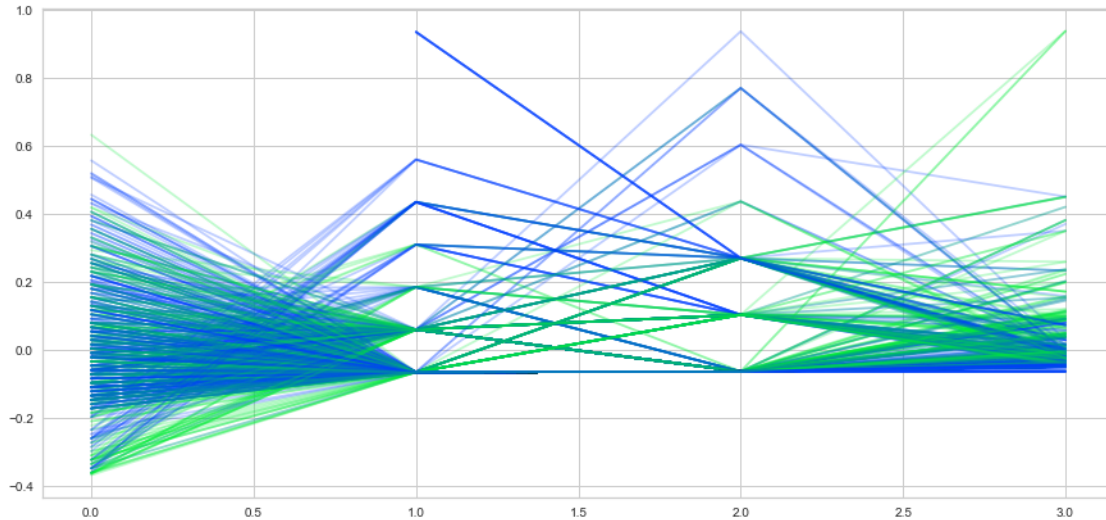
[14]: # setup the color for yellowbrick visualizer
      from yellowbrick.style import set_palette
      set_palette('sns_bright')
      # import packages
      from yellowbrick.features import ParallelCoordinates

[15]: # Specify the features of interest and the classes of the target
      classes = ['Not-survived', 'Survived']
      num_features = ['Age', 'SibSp', 'Parch', 'Fare']
      # copy data to a new dataframe
      data_norm = df_train.copy()

[16]: # normalize data to 0-1 range
      for feature in num_features:
          data_norm[feature] = (df_train[feature] - df_train[feature].
      ↪mean(skipna=True)) / (df_train[feature].max(skipna=True) - df_train[feature].
      ↪min(skipna=True))

[17]: # Extract the numpy arrays from the data frame
      X = data_norm[num_features].to_numpy()
      y = df_train.Survived.to_numpy()

      # Instantiate the visualizer
      visualizer = ParallelCoordinates(classes=classes, features=num_features)
      visualizer.fit(X, y)          # Fit the data to the visualizer
      visualizer.transform(X)      # Transform the data
      #visualizer.poof(outpath="d://pcoords2.png") # Draw/show/poof the data
      plt.show();
```



Passengers traveling with siblings on the boat have a higher death rate

passengers who paid a higher fare had a higher survival rate.

```
[18]: #10. Use Stack Bar Charts to compare passengers who survived to passengers who
      ↪ didn't survive based on the other variables.
      # a. More females survived than men. 3rd Class Tickets had a lower survival
      ↪ rate.
      # Also, Embarkation from Southampton port had a lower survival rate.
      plt.rcParams['figure.figsize'] = (20, 10)
      # make subplots
      fig, axes = plt.subplots(nrows = 2, ncols = 2)
      # make the data read to feed into the visualizer
      Sex_survived = df_train.replace({'Survived': {1: 'Survived', 0:
      ↪ 'Not-survived'}})[df_train['Survived']==1]['Sex'].value_counts()
      Sex_not_survived = df_train.replace({'Survived': {1: 'Survived', 0:
      ↪ 'Not-survived'}})[df_train['Survived']==0]['Sex'].value_counts()
      Sex_not_survived = Sex_not_survived.reindex(index = Sex_survived.index)
      # make the bar plot
      p1 = axes[0, 0].bar(Sex_survived.index, Sex_survived.values)
      p2 = axes[0, 0].bar(Sex_not_survived.index, Sex_not_survived.values,
      ↪ bottom=Sex_survived.values)
      axes[0, 0].set_title('Sex', fontsize=25)
      axes[0, 0].set_ylabel('Counts', fontsize=20)
      axes[0, 0].tick_params(axis='both', labelsize=15)
      axes[0, 0].legend((p1[0], p2[0]), ('Survived', 'Not-survived'), fontsize = 15)

      # make the data read to feed into the visualizer
```



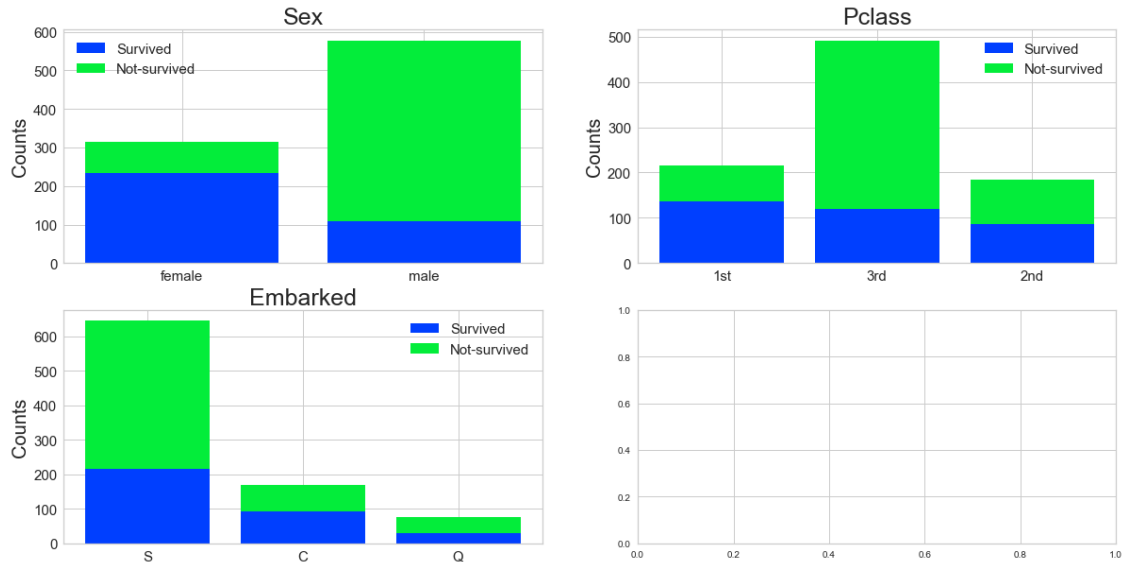
```

Pclass_survived = df_train.replace({'Survived': {1: 'Survived', 0:
↳ 'Not-survived'}}).replace({'Pclass': {1: '1st', 2: '2nd', 3:
↳ '3rd'}})[df_train['Survived']==1]['Pclass'].value_counts()
Pclass_not_survived = df_train.replace({'Survived': {1: 'Survived', 0:
↳ 'Not-survived'}}).replace({'Pclass': {1: '1st', 2: '2nd', 3:
↳ '3rd'}})[df_train['Survived']==0]['Pclass'].value_counts()
Pclass_not_survived = Pclass_not_survived.reindex(index = Pclass_survived.index)
# make the bar plot
p3 = axes[0, 1].bar(Pclass_survived.index, Pclass_survived.values)
p4 = axes[0, 1].bar(Pclass_not_survived.index, Pclass_not_survived.values,
↳ bottom=Pclass_survived.values)
axes[0, 1].set_title('Pclass', fontsize=25)
axes[0, 1].set_ylabel('Counts', fontsize=20)
axes[0, 1].tick_params(axis='both', labelsize=15)
axes[0, 1].legend((p3[0], p4[0]), ('Survived', 'Not-survived'), fontsize = 15)

# make the data read to feed into the visualizer
Embarked_survived = df_train.replace({'Survived': {1: 'Survived', 0:
↳ 'Not-survived'}})[df_train['Survived']==1]['Embarked'].value_counts()
Embarked_not_survived = df_train.replace({'Survived': {1: 'Survived', 0:
↳ 'Not-survived'}})[df_train['Survived']==0]['Embarked'].value_counts()
Embarked_not_survived = Embarked_not_survived.reindex(index = Embarked_survived.
↳ index)
# make the bar plot
p5 = axes[1, 0].bar(Embarked_survived.index, Embarked_survived.values)
p6 = axes[1, 0].bar(Embarked_not_survived.index, Embarked_not_survived.values,
↳ bottom=Embarked_survived.values)
axes[1, 0].set_title('Embarked', fontsize=25)
axes[1, 0].set_ylabel('Counts', fontsize=20)
axes[1, 0].tick_params(axis='both', labelsize=15)
axes[1, 0].legend((p5[0], p6[0]), ('Survived', 'Not-survived'), fontsize = 15)
plt.show()

```

[18]: <matplotlib.legend.Legend at 0x245f52426a0>



```
[19]: # 11. Some of my questions have been answered by seeing the charts but in some
      ↪ ways,
      # looking at this much data has created even more questions.
      # a. Now it's time to reduce some of the features so we can concentrate on the
      ↪ things that matter!
      # There features we will get rid of are: "PassengerId", "Name", "Ticket" and
      ↪ "Cabin".
      # (ID doesn't really give us any useful data, Ticket and Cabin have too many
      ↪ variables.
      # Name might reflect that they are related but we're keeping the category
      ↪ about siblings (for now).
```

```
[20]: # b. We can also fill in missing values.
      # (Cabin has some missing values but we are dropping that feature.)
      # Age has some missing values so I'll fill in with the average age.
      # Embarked also has some missing so I'll the most common.
```

```
[21]: # fill the missing age data with median value
def fill_na_median(df_train, inplace=True):
    return df_train.fillna(df_train.median(), inplace=inplace)

fill_na_median(df_train['Age'])

# check the result
print(df_train['Age'].describe())
```

```
count    891.000000
mean      29.361582
```

```

std      13.019697
min       0.420000
25%      22.000000
50%      28.000000
75%      35.000000
max      80.000000
Name: Age, dtype: float64

```

```

[22]: # fill with the most represented value
def fill_na_most(df_train, inplace=True):
    return df_train.fillna('S', inplace=inplace)

fill_na_most(df_train['Embarked'])

# check the result
print(df_train['Embarked'].describe())

```

```

count      891
unique       3
top         S
freq       646
Name: Embarked, dtype: object

```

```

[23]: # import package
import numpy as np

# log-transformation
def log_transformation(df_train):
    return df_train.apply(np.log1p)

df_train['Fare_log1p'] = log_transformation(df_train['Fare'])

# check the data
print(df_train.describe())

```

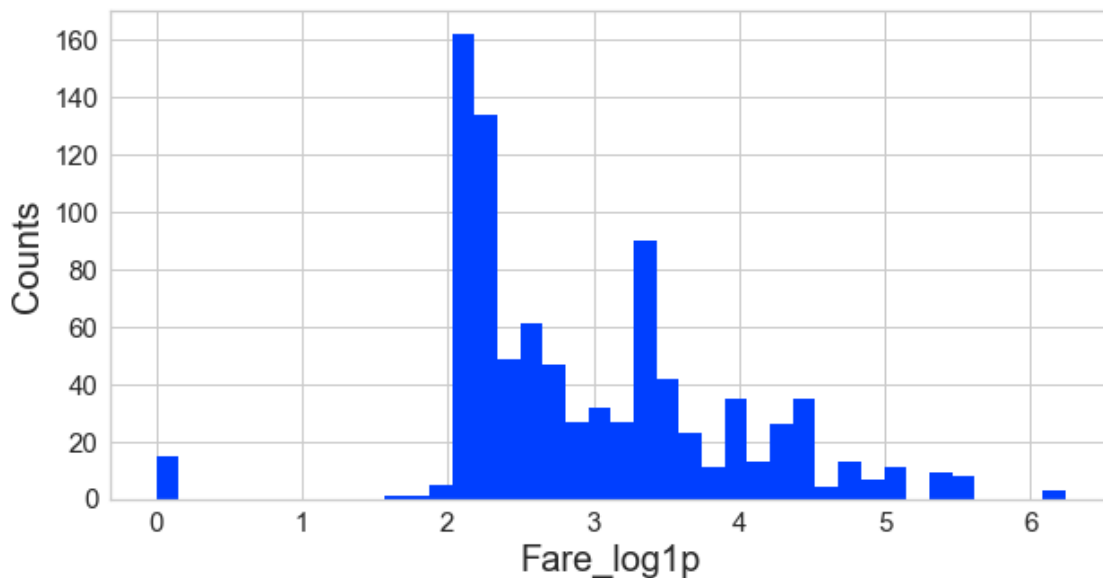
	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	891.000000	891.000000
mean	446.000000	0.383838	2.308642	29.361582	0.523008
std	257.353842	0.486592	0.836071	13.019697	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	22.000000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	35.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare	Fare_log1p
count	891.000000	891.000000	891.000000
mean	0.381594	32.204208	2.962246
std	0.806057	49.693429	0.969048

min	0.000000	0.000000	0.000000
25%	0.000000	7.910400	2.187218
50%	0.000000	14.454200	2.737881
75%	0.000000	31.000000	3.465736
max	6.000000	512.329200	6.240917

```
[24]: #12 If you go back and look at the histograms of Fare, you'll see that it is
      ↪very skewed...many low cost fares, not very many high cost fares.
      # Log Transformation is a good method to use on highly skewed data.
      #check the distribution using histogram
      # set up the figure size
      #%matplotlib inline
      plt.rcParams['figure.figsize'] = (10, 5)

      plt.hist(df_train['Fare_log1p'], bins=40)
      plt.xlabel('Fare_log1p', fontsize=20)
      plt.ylabel('Counts', fontsize=20)
      plt.tick_params(axis='both', labelsize=15)
      #plt.show()
```



```
[25]: # Convert your categorical data into numbers (Sex, PClass, Embark)
      #Step 13 - convert categorical data to numbers
      #get the categorical data
      cat_features = ['Pclass', 'Sex', "Embarked"]
      data_cat = df_train[cat_features]
```

```
[26]: data_cat.head()
```

```
[26]:   Pclass      Sex Embarked
      0         3   male        S
      1         1 female        C
      2         3 female        S
      3         1 female        S
      4         3   male        S
```

```
[27]: data_cat = data_cat.replace({'Pclass': {1: '1st', 2: '2nd', 3: '3rd'}})
```

```
[28]: data_cat.head()
```

```
[28]:   Pclass      Sex Embarked
      0     3rd    male        S
      1     1st  female        C
      2     3rd  female        S
      3     1st  female        S
      4     3rd    male        S
```

```
[29]: # One Hot Encoding
data_cat_dummies = pd.get_dummies(data_cat)
# check the data
print(data_cat_dummies.head(8))
```

	Pclass_1st	Pclass_2nd	Pclass_3rd	Sex_female	Sex_male	Embarked_C	\
0	0	0	1	0	1	0	
1	1	0	0	1	0	1	
2	0	0	1	1	0	0	
3	1	0	0	1	0	0	
4	0	0	1	0	1	0	
5	0	0	1	0	1	0	
6	1	0	0	0	1	0	
7	0	0	1	0	1	0	

	Embarked_Q	Embarked_S
0	0	1
1	0	0
2	0	1
3	0	1
4	0	1
5	1	0
6	0	1
7	0	1

CASE STUDY -3

```
[30]: #Step 14 - create a whole features dataset that can be used for train and
      ↪ validation data splitting
      # here we will combine the numerical features and the dummie features together
```

```
features_model = ['Age', 'SibSp', 'Parch', 'Fare_log1p']
data_model_X = pd.concat([df_train[features_model], data_cat_dummies], axis=1)
```

```
[31]: data_model_X.head()
```

```
[31]:
```

	Age	SibSp	Parch	Fare_log1p	Pclass_1st	Pclass_2nd	Pclass_3rd	\
0	22.0	1	0	2.110213	0	0	1	
1	38.0	1	0	4.280593	1	0	0	
2	26.0	0	0	2.188856	0	0	1	
3	35.0	1	0	3.990834	1	0	0	
4	35.0	0	0	2.202765	0	0	1	

	Sex_female	Sex_male	Embarked_C	Embarked_Q	Embarked_S
0	0	1	0	0	1
1	1	0	1	0	0
2	1	0	0	0	1
3	1	0	0	0	1
4	0	1	0	0	1

```
[32]: # create a whole target dataset that can be used for train and validation data
      ↪ splitting
      #data_model_y1 = df_train.replace({'Survived': {1: 'Survived', 0:
      ↪ 'Not_survived'}})['Survived']
      data_model_y = df_train['Survived']
      data_model_y.head()
```

```
[32]: 0    0
      1    1
      2    1
      3    1
      4    0
      Name: Survived, dtype: int64
```

```
[34]: # separate data into training and validation and check the details of the
      ↪ datasets
      # import packages
      from sklearn.model_selection import train_test_split

      # split the data
      X_train, X_val, y_train, y_val = train_test_split(data_model_X, data_model_y,
      ↪ test_size =0.3, random_state=11)

      # number of samples in each set
      print("No. of samples in training set: ", X_train.shape[0])
      print("No. of samples in validation set:", X_val.shape[0])

      # Survived and not-survived
```

```

print('\n')
print('No. of survived and not-survived in the training set:')
print(y_train.value_counts())

print('\n')
print('No. of survived and not-survived in the validation set:')
print(y_val.value_counts())

```

No. of samples in training set: 623
 No. of samples in validation set: 268

No. of survived and not-survived in the training set:
 0 373
 1 250
 Name: Survived, dtype: int64

No. of survived and not-survived in the validation set:
 0 176
 1 92
 Name: Survived, dtype: int64

```

[38]: # Step 15 - Eval Metrics
from sklearn.linear_model import LogisticRegression

from yellowbrick.classifier import ConfusionMatrix
from yellowbrick.classifier import ClassificationReport
from yellowbrick.classifier import ROCAUC
# ROC and AUC
# Instantiate the classification model
model = LogisticRegression(solver='liblinear')
#The ConfusionMatrix visualizer takes a model
classes = ['Not_survived', 'Survived']
cm = ConfusionMatrix(model, classes=classes, label_encoder={0: "Not_survived",
    ↳1: "Survived"}, percent=False)

#Fit fits the passed model. This is unnecessary if you pass the visualizer a
    ↳pre-fitted model
cm.fit(X_train, y_train)

#To create the ConfusionMatrix, we need some test data. Score runs predict() on
    ↳the data
#and then creates the confusion_matrix from scikit learn.
cm.score(X_val, y_val)

# change fontsize of the labels in the figure


```

```
for label in cm.ax.texts:
    label.set_size(20)
```

```
#How did we do?
cm.poof()
```

C:\ProgramData\Anaconda3\lib\site-packages\sklearn\metrics_classification.py:191: FutureWarning: elementwise comparison failed; returning scalar instead, but in the future will perform elementwise comparison


```
>>> y_true = [0, 1, 2, 3]
```




```
[39]: # Step 15 - Eval Metrics
# Precision, Recall, and F1 Score
# set the size of the figure and the font size
##matplotlib inline
plt.rcParams['figure.figsize'] = (15, 7)
plt.rcParams['font.size'] = 20

# Instantiate the visualizer
visualizer = ClassificationReport(model, classes=classes)

visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_val, y_val) # Evaluate the model on the test data
g = visualizer.poof()
```

```
[40]: # ROC and AUC
#Instantiate the visualizer
visualizer = ROCAUC(model, pos_label = 1)
visualizer.fit(X_train, y_train) # Fit the training data to the visualizer
visualizer.score(X_val, y_val) # Evaluate the model on the test data
g = visualizer.poof()
```



```
[ ]:
```