

Assignment_9_3_Raghuwanshi_Prashant_DSC550

November 4, 2021

Assignment: 9.3 Exercise: Neural Network Classifiers

Name: Prashant Raghuwanshi

Date: 10/29/2021

Course: DSC550-T301 Data Mining (2221-1) Neural Network Classifier with Scikit

Using the multi-label classifier dataset (categorized-comments.jsonl), fit a neural network classifier using scikit-learn to predict the comment category. Use the code found in chapter 12 of the Applied Text Analysis with Python book as a guide, but you will need to modify the code for this dataset. Report the accuracy, precision, recall, F1-score, and confusion matrix.

```
[1]: import pandas as pd
import numpy as np
import jsonlines
import time
from textblob import TextBlob
import json

[2]: # 1.Load the data from the "categorized-comments.jsonl" file into a DataFrame.
addr1 = "C:/Users/dell/Documents/Machine_learning_assignments/week-9/
↪categorized-comments.jsonl"

[3]: # Reading the JSON line file into a dataframe
newlist = []
with jsonlines.open(addr1) as f:
    for obj in f.iter(type=dict, skip_invalid=True):
        newlist.append(obj)
ccjsonsrc = pd.DataFrame(newlist)
ccjsonsrc.head(5)
```

	cat	txt
0	sports	Barely better than Gabbert? He was significant...
1	sports	Fuck the ducks and the Angels! But welcome to ...
2	sports	Should have drafted more WRs.\n\n- Matt Millen...
3	sports	[Done](https://i.imgur.com/2YZ90pm.jpg)
4	sports	No!! NOO!!!!

```
[4]: # using lambda function and convert the string to lower case
ccjsonlower = ccjsonsrc.apply(lambda x: x.astype(str).str.lower())
# limiting the records in dataframe
sampledf = ccjsonlower.head(80000)
sampledf.head()
```

```
[4]:      cat      txt
0  sports  barely better than gabbert? he was significant...
1  sports  fuck the ducks and the angels! but welcome to ...
2  sports  should have drafted more wrs.\n\n- matt millen...
3  sports      [done](https://i.imgur.com/2yz90pm.jpg)
4  sports      no!! noo!!!!
```

```
[5]: #Remove all punctuation from the text.
# Create the punctuation dictionary by using unicodedata
import sys
import unicodedata
punctuation = dict.fromkeys(i for i in range(sys.maxunicode)
                             if unicodedata.category(chr(i)).startswith('P'))
```

```
[6]: # removing punctuation from each row of dataframe's txt column
for i in range(len(sampledf)):
    test = [string.translate(punctuation) for string in (sampledf.loc[i,
↪ "txt"])]
    # converting list to string
    test1 = "".join(str(x) for x in test)
    # updating the row values
    sampledf.loc[i, ["txt"]] = test1
# print dataframe after removing punctuations from txt column
sampledf.head()
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1637:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self._setitem_single_block(indexer, value, name)

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:692:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

iloc._setitem_with_indexer(indexer, value, self.name)

```
[6]:      cat      txt
0  sports  barely better than gabbert he was significantl...
```

```

1 sports  fuck the ducks and the angels but welcome to a...
2 sports  should have drafted more wrs\n\n matt millen p...
3 sports                                     donehttpsiimgurcom2yz90pmjpg
4 sports                                     no noo

```

```

[7]: # Remove stop words.
      # load library
      import nltk
      from nltk.corpus import stopwords
      from nltk.tokenize import word_tokenize
      # load stop words
      stop_words = stopwords.words('english')

```

```

[8]: # remove stop words from each row of dataframe's txt column
      for i in range(len(sampledf)) :
          # tokenized each row of dataframe's txt column
          test_token = word_tokenize(sampledf.loc[i, "txt"])
          # remove stop words
          rem_words = [word for word in test_token if word not in stop_words]
          # coverting list to string
          rem_words1 = " ".join(str(x) for x in rem_words)
          # writting back processed removed stop words to dataframe
          # updating the row values for txt column
          sampledf.loc[i, ["txt"]] = rem_words1
      # printing last rows of dataframe showing removed stop words
      print(rem_words)

```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1637:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self._setitem_single_block(indexer, value, name)

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:692:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

iloc._setitem_with_indexer(indexer, value, self.name)

['love', 'latest', 'deck', 'video']

```

[9]: #print dataframe after updating the removed stop words to txt column
      sampledf.head()

```

```
[9]:      cat                                txt
0  sports  barely better gabbert significantly better yea...
1  sports                fuck ducks angels welcome new niners fans
2  sports                drafted wrs matt millen probably
3  sports                donehttpsiimgurcom2yz90pmjpg
4  sports                noo
```

```
[10]: # load library
from nltk.stem.porter import PorterStemmer
# create stemmer
porter = PorterStemmer()
# apply stemmer to each row of dataframe's txt column
for i in range(len(sampledf)) :
    # tokenized each row of dataframe's txt column
    test_token1 = word_tokenize(sampledf.loc[i, "txt"])
    # apply stemmer
    porter_words = [porter.stem(word) for word in test_token1]
    # coverting list to string
    porter_words1 = " ".join(str(x) for x in porter_words)
    # writting back processed removed stop words to dataframe
    # updating the row values for txt column
    sampledf.loc[i, ["txt"]] = porter_words1
# printing last rows of dataframe showing removed stop words
print(porter_words)
```

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:1637:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

self._setitem_single_block(indexer, value, name)

C:\ProgramData\Anaconda3\lib\site-packages\pandas\core\indexing.py:692:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

iloc._setitem_with_indexer(indexer, value, self.name)

['love', 'latest', 'deck', 'video']

```
[11]: # print dataframe after updating the applied stemmer to each row of dataframe's
      ↪txt column
sampledf.head()
```

```
[11]:      cat                                txt
0  sports  bare better gabbert significantli better year ...
```

```

1  sports          fuck duck angel welcom new niner fan
2  sports          draft wr matt millen probabl
3  sports          donehttpsiimgurcom2yz90pmjpg
4  sports          noo

```

```
[12]: sampledf3 = sampledf
```

```
[13]: from io import StringIO
sampledf3['category_id'] = sampledf3['cat'].factorize()[0]
category_id_df = sampledf3[['cat', 'category_id']].drop_duplicates().
    ↪sort_values('category_id')
category_to_id = dict(category_id_df.values)
id_to_category = dict(category_id_df[['category_id', 'cat']].values)
sampledf3.head()
```


<ipython-input-13-8e3e7fec068a>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy
sampledf3['category_id'] = sampledf3['cat'].factorize()[0]

```
[13]:      cat          txt  category_id
0  sports  bare better gabbert significantli better year ...      0
1  sports          fuck duck angel welcom new niner fan      0
2  sports          draft wr matt millen probabl      0
3  sports          donehttpsiimgurcom2yz90pmjpg      0
4  sports          noo      0

```

```
[14]: # verifying Imbalanced Classes
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,6))
sampledf3.groupby('cat').txt.count().plot.bar(ylim=0)
plt.show()
```



```
[15]: print(sampledf3.columns)
```

```
Index(['cat', 'txt', 'category_id'], dtype='object')
```

```
[16]: sampledf4 = sampledf3.head(40000)
sampledf5 = sampledf4
```

```
[17]: # Convert each entry into a term frequency-inverse document frequency (tfidf)
      ↪ vector
      #### Weighted Word Importance
      #### here we are comparing the frequency of words in a document (a tweet,
      ↪ movie review speech transcript)
      #### with the frequency of words in all other documents using term
      ↪ frequency-inverse document frequency
```

```
# import libraries
from sklearn.feature_extraction.text import TfidfVectorizer
ftidf = TfidfVectorizer()
```

```
[18]: from sklearn.feature_extraction.text import TfidfVectorizer
      tfidf = TfidfVectorizer(sublinear_tf=True, min_df=5, norm='l2',
      ↪encoding='latin-1', ngram_range=(1, 2), stop_words='english')
      features = tfidf.fit_transform(sampledf4.txt).toarray()
      labels = sampledf4.category_id
      features.shape
```

```
[18]: (40000, 15229)
```

```
[19]: from sklearn.model_selection import train_test_split
      from sklearn.feature_extraction.text import CountVectorizer
      from sklearn.feature_extraction.text import TfidfTransformer
      from sklearn.naive_bayes import MultinomialNB
      X_train, X_test, y_train, y_test = train_test_split(sampledf4['txt'],
      ↪sampledf4['category_id'], random_state = 0)
      count_vect = CountVectorizer()
      X_train_counts = count_vect.fit_transform(X_train)
      tfidf_transformer = TfidfTransformer()
      X_train_tfidf = tfidf_transformer.fit_transform(X_train_counts).toarray()
      #clf = MultinomialNB().fit(X_train_tfidf, y_train)
```

```
[20]: print(X_train.shape); print(X_test.shape); print(y_train.shape); print(y_test.
      ↪shape);print(X_train_tfidf.shape)
```

```
(30000,)
(10000,)
(30000,)
(10000,)
(30000, 24587)
```

```
[21]: from sklearn.neural_network import MLPClassifier

      mlp = MLPClassifier(hidden_layer_sizes=(500,500,500), activation='relu',
      ↪solver='adam', max_iter=5)
      #mlp = MLPClassifier(hidden_layer_sizes=(500, 150))
      #mlp.fit(X_train,y_train)
      mlp.fit(X_train_tfidf, y_train)
```

```
C:\ProgramData\Anaconda3\lib\site-
packages\sklearn\neural_network\_multilayer_perceptron.py:582:
ConvergenceWarning: Stochastic Optimizer: Maximum iterations (5) reached and the
optimization hasn't converged yet.
      warnings.warn(
```

```
[21]: MLPClassifier(hidden_layer_sizes=(500, 500, 500), max_iter=5)
```

```
[22]: predict_train = mlp.predict(X_train_tfidf)
```

```
[23]: ## we can evaluate the performance of the model. Being a classification  
→algorithm,  
# we will first import the required modules, which is done in the first line of  
→code below.  
# The second and third lines of code print the confusion matrix and the  
→confusion report results on the training data.  
from sklearn.metrics import classification_report, confusion_matrix  
print(confusion_matrix(y_train, predict_train))  
print(classification_report(y_train, predict_train))  
#print(classification_report(y_train, predict_train))  
#print(confusion_matrix(y_train.argmax(axis=1), predict_train.argmax(axis=1)))
```

```
[[10433  739]
 [   51 18777]]
```

	precision	recall	f1-score	support
0	1.00	0.93	0.96	11172
1	0.96	1.00	0.98	18828
accuracy			0.97	30000
macro avg	0.98	0.97	0.97	30000
weighted avg	0.97	0.97	0.97	30000

```
[24]: ## The above output shows the performance of the model on training data. The  
→accuracy and the F1 score is around 0.98 and 0.99
```

```
[25]: X_test_counts = count_vect.fit_transform(X_test)  
tfidf_transformer = TfidfTransformer()  
X_test_tfidf = tfidf_transformer.fit_transform(X_test_counts).toarray()
```

```
[26]: X_test_tfidf.shape
```

```
[26]: (10000, 13354)
```

```
[27]: X_test1_tfidf = np.resize(X_test_tfidf, (10000, 24587))
```

```
[28]: X_test1_tfidf.shape
```

```
[28]: (10000, 24587)
```

```
[29]: predict_test = mlp.predict(X_test1_tfidf)
```



```
[30]: print(confusion_matrix(y_test, predict_test))
      print(classification_report(y_test, predict_test))
```

```
[[1688 2029]
 [2858 3425]]
```

	precision	recall	f1-score	support
0	0.37	0.45	0.41	3717
1	0.63	0.55	0.58	6283
accuracy			0.51	10000
macro avg	0.50	0.50	0.50	10000
weighted avg	0.53	0.51	0.52	10000

```
[31]: ## The above output shows the performance of the model on test data. The
      → accuracy and the F1 score is around 0.54 and 0.53
```

0.0.1 Neural Network Classifier with Keras

```
[32]: # Keras specific
import keras
from keras.models import Sequential
from keras.layers import Dense
from tensorflow.keras.utils import to_categorical
```

```
[33]: # Since our target variable represents a binary category which has been coded
      → as numbers 0 and 1, we will have to encode it.
      # We can easily achieve that using the "to_categorical" function from the Keras
      → utilities package.
      # The two lines of code below accomplishes that in both training and test
      → datasets.
      # one hot encode outputs
y_train2 = to_categorical(y_train)
y_test2 = to_categorical(y_test)

count_classes = y_test2.shape[1]
print(count_classes)
```

2

```
[34]: print(X_train_tfidf.shape, y_train2.shape)
```

```
(30000, 24587) (30000, 2)
```

```
[47]: # downsizing the records due to memory issue
X_train5 = X_train.head(100)
```

```
[41]: # converting words to numbers
count = CountVectorizer()
bag_of_words = count.fit_transform(X_train5.to_numpy())
X_train4 = bag_of_words.toarray()
```

```
[43]: print(X_train4.shape, y_train2.shape)
```

```
(100, 747) (30000, 2)
```

```
[59]: y_train4 = np.resize(y_train2, (100, 2))
```

```
[60]: y_train4.shape
```

```
[60]: (100, 2)
```

```
[53]: model = Sequential()
model.add(Dense(500, activation='relu', input_dim=747))
model.add(Dense(100, activation='relu'))
model.add(Dense(50, activation='relu'))
model.add(Dense(2, activation='softmax'))

# Compile the model
model.compile(loss='categorical_crossentropy', optimizer='adam',
              metrics=['accuracy'])
```

```
[61]: import tensorflow as tf
# build the model
tf.config.run_functions_eagerly(True)
#model.fit(X_train, y_train, epochs=20)
model.fit(X_train4, y_train4, epochs=20)
```

```
Epoch 1/20
```

```
4/4 [=====] - 1s 43ms/step - loss: 0.6917 - accuracy: 0.5200
```

```
Epoch 2/20
```

```
4/4 [=====] - 0s 31ms/step - loss: 0.5592 - accuracy: 0.7200
```

```
Epoch 3/20
```

```
4/4 [=====] - 0s 39ms/step - loss: 0.4653 - accuracy: 0.7900
```

```
Epoch 4/20
```

```
4/4 [=====] - 0s 42ms/step - loss: 0.3845 - accuracy: 0.9200
```

```
Epoch 5/20
```

```
4/4 [=====] - 0s 37ms/step - loss: 0.3066 - accuracy: 0.9700
```

```
Epoch 6/20
```

```
4/4 [=====] - 0s 31ms/step - loss: 0.2354 - accuracy: 0.9800
```

```

Epoch 7/20
4/4 [=====] - 0s 28ms/step - loss: 0.1711 - accuracy:
0.9900
Epoch 8/20
4/4 [=====] - 0s 29ms/step - loss: 0.1218 - accuracy:
0.9900
Epoch 9/20
4/4 [=====] - 0s 31ms/step - loss: 0.0862 - accuracy:
0.9900
Epoch 10/20
4/4 [=====] - 0s 31ms/step - loss: 0.0614 - accuracy:
0.9900
Epoch 11/20
4/4 [=====] - 0s 33ms/step - loss: 0.0454 - accuracy:
0.9900
Epoch 12/20
4/4 [=====] - 0s 40ms/step - loss: 0.0360 - accuracy:
0.9900
Epoch 13/20
4/4 [=====] - 0s 40ms/step - loss: 0.0315 - accuracy:
0.9900
Epoch 14/20
4/4 [=====] - 0s 33ms/step - loss: 0.0251 - accuracy:
0.9900
Epoch 15/20
4/4 [=====] - 0s 33ms/step - loss: 0.0316 - accuracy:
0.9800
Epoch 16/20
4/4 [=====] - ETA: 0s - loss: 0.0330 - accuracy: 0.98 -
0s 48ms/step - loss: 0.0330 - accuracy: 0.9800
Epoch 17/20
4/4 [=====] - 0s 34ms/step - loss: 0.0311 - accuracy:
0.9800
Epoch 18/20
4/4 [=====] - 0s 29ms/step - loss: 0.0285 - accuracy:
0.9800
Epoch 19/20
4/4 [=====] - 0s 46ms/step - loss: 0.0248 - accuracy:
0.9700
Epoch 20/20
4/4 [=====] - 0s 35ms/step - loss: 0.0217 - accuracy:
0.9900

```

[61]: <keras.callbacks.History at 0x17cf23997f0>

```

[62]: from sklearn.metrics import f1_score, precision_score, recall_score,
      ↪ confusion_matrix

```

```
y_pred1 = model.predict(X_train4)
y_pred = np.argmax(y_pred1, axis=1)
```

```
[63]: y_train6 = np.resize(y_train2, (100))
```

```
[64]: from sklearn.metrics import classification_report, confusion_matrix
print(confusion_matrix(y_train6, y_pred))
print(classification_report(y_train6, y_pred))
```

```
[[16 34]
 [19 31]]
```

	precision	recall	f1-score	support
0.0	0.46	0.32	0.38	50
1.0	0.48	0.62	0.54	50
accuracy			0.47	100
macro avg	0.47	0.47	0.46	100
weighted avg	0.47	0.47	0.46	100

```
[ ]: ## The above output shows the performance of the model on train data. The
      ↪ accuracy and the F1 score is around 0.47 and 0.46
```

```
[65]: X_test.shape
```

```
[65]: (10000,)
```

```
[66]: # downsizing the test records, due to memory issue
X_test3 = X_test.head(100)
```

```
[67]: # converting words to numbers
bag_of_words1 = count.fit_transform(X_test3.to_numpy())
X_test4 = bag_of_words1.toarray()
```

```
[68]: X_test4.shape
```

```
[68]: (100, 811)
```

```
[69]: # resizing the array, due to difference in feature of trained model dataset and
      ↪ test model data set
X_test6 = np.resize(X_test4, (100, 747))
```

```
[70]: # evaluating the model with test data
y_test_pred1 = model.predict(X_test6)
y_test_pred = np.argmax(y_pred1, axis=1)
```

```
[71]: # resizing the array, due to difference in counts of test model dataset and
      ↪ test model data set
      y_test3 = np.resize(y_test, (100))
```

```
[72]: print(confusion_matrix(y_test3,y_test_pred))
      print(classification_report(y_test3,y_test_pred))
```

```
[[12 27]
 [23 38]]
```

	precision	recall	f1-score	support
0	0.34	0.31	0.32	39
1	0.58	0.62	0.60	61
accuracy			0.50	100
macro avg	0.46	0.47	0.46	100
weighted avg	0.49	0.50	0.49	100

```
[74]: ## The above output shows the performance of the model on test data. The
      ↪ accuracy and the F1 score is around 0.49 and 0.49
```

Classifying Images

```
[75]: from keras.datasets import mnist
      from keras.models import sequential
      from keras.layers import Dense, Dropout, Flatten
      from keras.layers.convolutional import Conv2D, MaxPooling2D
      from keras.utils import np_utils
      from keras import backend as k
```

```
[76]: # ste that the color channel value will be first
      k.set_image_data_format("channels_last")
```

```
[77]: #set seed
      np.random.seed(0)
```

```
[78]: # set image information
      channels = 1
      height = 28
      width =28
```

```
[79]: # load data and target from MNIST data
      (data_train, target_train), (data_test, target_test) = mnist.load_data()
```

```
[80]: # reshape training image data into features
      data_train = data_train.reshape(data_train.shape[0], height, width, channels)
```

```

[81]: # reshape test image data into features
data_test = data_test.reshape(data_test.shape[0], height, width, channels)

[83]: # rescale pixel intensity to between 0 and 1
features_train = data_train/255
features_test = data_test/255

[84]: # one hot encoding target
target_train = np_utils.to_categorical(target_train)
target_test = np_utils.to_categorical(target_test)
number_of_classes = target_test.shape[1]

[85]: # start neural network
network = Sequential()

[86]: # add convolutional layer with 64 filters
network.add(Conv2D(filters = 64, kernel_size=(5, 5), input_shape=(height, width, channels), activation='relu'))

#network.add(Conv2D(filters = 64, kernel_size=(5, 5), input_shape= tf.
#reshape(2, (1, 2, 1)), activation='relu'))

[87]: # add max pooling layer
network.add(MaxPooling2D(pool_size=(2,2), data_format='channels_last'))

[88]: # add dropout layer
network.add(Dropout(0.5))

[89]: network.add(Flatten())

[90]: network.add(Dense(128, activation="relu"))

[91]: #add dropout layer
network.add(Dropout(0.5))

[92]: # Add dropout layer
network.add(Dense(number_of_classes, activation="softmax"))

[93]: # compile neural network\
network.compile(loss="categorical_crossentropy", optimizer = "rmsprop",
metrics=["accuracy"])

[94]: print(features_train.shape); print(target_train.shape); print(features_test.
shape), print(target_test.shape)

(60000, 28, 28, 1)
(60000, 10)

```

```
(10000, 28, 28, 1)
(10000, 10)
```

```
[94]: (None, None)
```

```
[95]: # train neural network
network.fit(features_train, target_train, epochs=2, verbose=0, batch_size=1000,
↳ validation_data=(features_test, target_test))
```

```
[95]: <keras.callbacks.History at 0x17cf1fab5e0>
```

```
[ ]:
```